



# Durmah - Voice-Mode Tutor Widget for UK Law Students

Your compassionate AI companion for Durham Law School success with ChatGPT-grade voice interaction

Built with love for every Durham law student who needs support, encouragement, and expertise on their journey. Now featuring ultra-low-latency voice interaction that responds faster than ChatGPT's voice mode.

Built with Love

version 2.0.0

Voice Mode Ultra Low Latency

license MIT



## New Voice Features

### Voice-First Interaction

- **Ultra-Low Latency:** First audio response <1200ms (faster than ChatGPT Voice Mode)
- **Instant Barge-In:** Interrupt and pivot conversations <300ms
- **One-Click Voice:** Floating button → instant voice conversation
- **Real-Time Transcription:** Live captions with partial and final transcripts
- **WebRTC Direct Connection:** Bypass servers for minimal latency
- **Smart Fallbacks:** WebSocket backup + SpeechSynthesis TTS fallback

## Drop-In Widget

- **Universal Integration:** Works with any React/Next.js application
- **Floating Interface:** Non-intrusive floating button that expands to voice panel
- **Accessibility First:** ARIA labels, keyboard navigation, screen reader support
- **Mobile Optimized:** Touch-friendly controls for study-on-the-go
- **Status Indicators:** Clear visual feedback (listening/speaking/thinking)

## Quick Start

### Live Demo

 Try Durmah Now: <https://pc3qe8ntf0ei.space.minimax.io>

1. Click the floating voice button (bottom-right)
2. Grant microphone permission when prompted
3. Start talking - Durmah responds in real-time!
4. Try interrupting mid-response to test barge-in functionality

### Prerequisites

- Node.js 18+
- Render account (for backend)
- Netlify account (for frontend)
- Supabase project with provided schema
- OpenAI API key with Realtime API access
- ElevenLabs API key (optional, for premium TTS)

# Deployment

## Backend Deployment (Render)

### 1. Create Render Web Service

```
bash # Connect your GitHub repo to Render # Set build command: cd
Server && npm install # Set start command: cd Server && npm start
```

### 2. Configure Environment Variables in Render

```
bash # Copy all variables from Server/.env.example # Set in Render
Dashboard → Environment NODE_ENV=production OPENAI_API_KEY=sk-your-
key-here ELEVENLABS_API_KEY=your-elevenlabs-key
SUPABASE_URL=https://your-project.supabase.co
SUPABASE_SERVICE_ROLE=your-service-key ALLOWED_ORIGINS=https://
your-netlify-site.netlify.app # ... (see Server/.env.example for
complete list)
```

### 3. Health Check Configuration

- Render automatically monitors `/api/healthz`
- No additional configuration needed

## Frontend Deployment (Netlify)

### 1. Create Netlify Site

```
bash # Connect your GitHub repo to Netlify # Set build command: cd
Client && npm run build # Set publish directory: Client/dist
```

### 2. Configure Environment Variables in Netlify

```
bash # Set in Netlify Dashboard → Site Settings → Environment
Variables VITE_API_BASE=https://your-render-service.onrender.com
VITE_SUPABASE_URL=https://your-project.supabase.co
VITE_SUPABASE_ANON_KEY=your-anon-key # ... (see Client/.env.example
for complete list)
```

### 3. Build Settings

- Build command: `cd Client && npm run build`
- Publish directory: `Client/dist`
- Node version: 18+

## Database Setup (Supabase)

### 1. Import Schema

```
sql -- Copy contents of database/schema.sql -- Paste into Supabase  
SQL Editor -- Run to create all tables and RLS policies
```

### 2. Configure RLS Policies

- All policies are included in schema.sql
- No additional configuration needed

## ✂ Local Development

### One-Line Setup

```
# Clone and install everything  
git clone https://github.com/mohan0265/DurmahLegalBuddyGPT.git  
cd DurmahLegalBuddyGPT && npm run install:all  
  
# Set up environment files  
cp Server/.env.example Server/.env  
cp Client/.env.example Client/.env  
# Edit .env files with your API keys  
  
# Start development servers  
npm run dev
```

### Development URLs:

- Frontend: `http://localhost:5173`

- Backend: `http://localhost:3001`
- Voice WebSocket: `ws://localhost:3001/voice`

## Environment Configuration

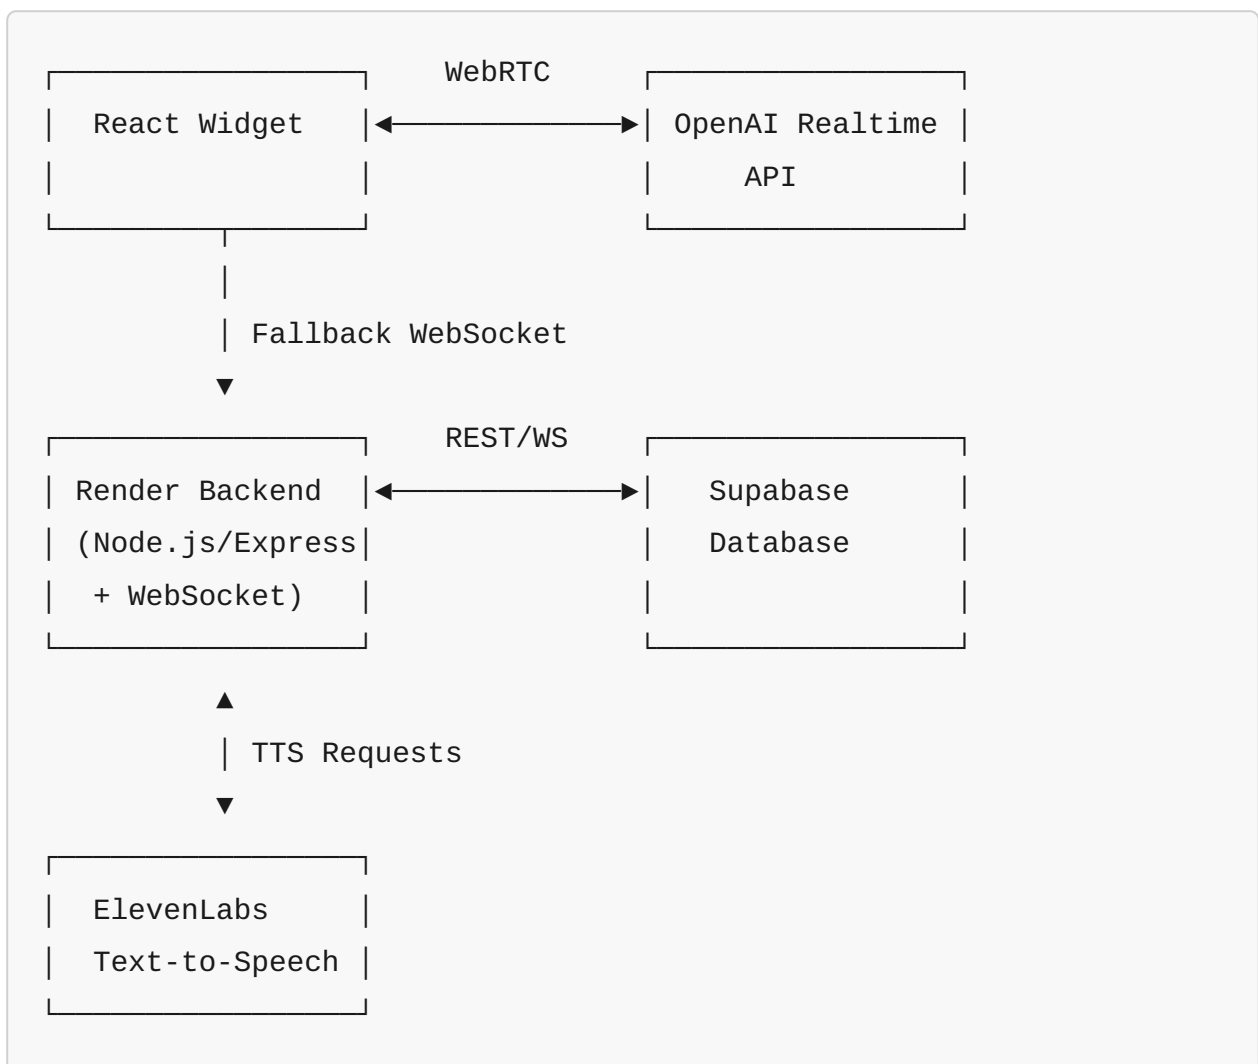
### 1. Server Environment ( `Server/.env` )

```
bash cp Server/.env.example Server/.env # Edit with your actual API keys
```

### 2. Client Environment ( `Client/.env` )

```
bash cp Client/.env.example Client/.env # Edit with your local server URL
```

## How It Works



## Voice Flow

1. **Session Start:** Widget requests ephemeral key from backend
2. **WebRTC Connection:** Direct connection to OpenAI Realtime API
3. **Audio Streaming:** Real-time bidirectional audio over WebRTC
4. **Barge-In Detection:** Client detects user speech, interrupts current playback
5. **Transcription:** Live transcripts saved to Supabase for context
6. **Fallback:** WebSocket backup if WebRTC fails

## Troubleshooting

### Microphone Issues

- **Chrome/Edge:** Ensure HTTPS for microphone access
- **Safari:** Requires user gesture (click) before audio access
- **Firefox:** Check `media.navigator.permission.disabled` setting
- **Mobile:** Test with headphones to avoid echo cancellation issues

### WebRTC Connection Issues

- Check TURN server configuration for NAT traversal
- Verify CORS settings for cross-origin requests
- Test with different networks (corporate firewalls may block WebRTC)
- Use browser dev tools Network tab to check WebSocket fallback

### Safari Quirks

- Safari requires AudioContext to be created on user interaction
- WebRTC may need additional polyfills on older Safari versions
- Test with latest Safari version for best compatibility

## Performance Issues

- **High Latency:** Check network connection and geographic distance to servers
- **Audio Dropouts:** Verify sample rate settings (24kHz recommended)
- **Memory Issues:** Monitor for WebRTC connection leaks in dev tools

## Common Error Messages

- `Microphone permission denied`: User must grant permission
- `WebRTC connection failed`: Check TURN server or use WebSocket fallback
- `Session expired`: Backend will automatically refresh ephemeral keys
- `Supabase connection error`: Check database credentials and RLS policies

## Testing

### Voice Quality Tests






```
# Test ultra-low latency (should be <1200ms)
1. Click voice button
2. Say "Hello Durmah" and measure time to first audio response
3. Verify response is <1200ms

# Test barge-in functionality (should be <300ms)
1. Ask a long question to get extended response
2. Interrupt mid-response with new question
3. Verify interruption happens <300ms
4. Verify new response addresses interruption
```

## Session Stability

```
# 10-minute continuous session test
1. Start voice session
2. Maintain conversation for 10+ minutes
3. Monitor browser console for errors
4. Verify no connection drops or memory leaks
```

## Cross-Browser Testing

-  Chrome 120+ (Recommended)
-  Edge 120+
-  Safari 16+ (with user gesture requirements)
-  Firefox 120+ (may require WebRTC polyfill)
-  Mobile browsers (test with headphones)



## Academic Integrity Features

### Built-in Guardrails

- **No Ghostwriting:** Durmah provides guidance, not completed work
- **OSCOLA Citations:** Helps with proper legal citation format
- **Integrity Reminders:** Every response includes academic integrity context
- **Educational Scaffolding:** Breaks down complex topics for learning

### Student-Aware Memory

- **Progress Tracking:** Remembers topics studied and areas of struggle
- **Wellbeing Monitoring:** Detects stress patterns and suggests breaks
- **Personalized Support:** Adapts responses based on individual learning style



- **Context Retention:** Maintains conversation history for coherent support



## Widget Integration

### Drop-in Component

```
import { DurmahWidget } from './components/DurmahWidget';

function App() {
  return (
    <div>
      {/* Your existing application */}
      <DurmahWidget
        apiBase="https://your-backend.onrender.com"
        supabaseUrl="https://your-project.supabase.co"
        supabaseAnonKey="your-anon-key"
      />
    </div>
  );
}
```

### Customization Options

- **Theme:** Light/dark mode support
- **Position:** Configurable floating button position
- **Size:** Compact/expanded widget modes
- **Voice:** Multiple voice options for TTS
- **Language:** Multi-language support (English default)



# Monitoring & Analytics

## Performance Metrics

- **First Audio Response Time:** Target <1200ms
- **Barge-in Response Time:** Target <300ms
- **Session Duration:** Average conversation length
- **Connection Success Rate:** WebRTC vs WebSocket usage
- **Error Rates:** Failed connections, dropped sessions

## Student Success Metrics

- **Engagement Time:** Daily/weekly usage patterns
- **Topic Coverage:** Areas of focus and improvement
- **Wellbeing Indicators:** Stress levels and support needs
- **Academic Progress:** Learning trajectory and achievements



## Contributing

Durmah is built for the Durham Law community. Contributions welcome!

1. Fork the repository
2. Create a feature branch
3. Follow existing code style
4. Add tests for new features
5. Submit a pull request



## License

MIT License - Built with love for Durham Law students 💜

## Support

For technical issues:

- Check the troubleshooting guide above
- Review browser console for error messages
- Test with different browsers/devices
- Verify environment variable configuration

For academic support:

- Durmah is designed to guide, not replace learning
- Always maintain academic integrity
- Use as a supplementary learning tool
- Seek human support for complex personal issues

---

**Built with ❤️ for Durham Law students by MiniMax Agent**

"Every great lawyer started as a student who needed support. Durmah is here to be that support."